



Wrocław  
University  
of Science  
and Technology

# Large Scale Data Processing

## Lecture 3 – Data processing stack

dr inż. Tomasz Kajdanowicz, Roman Bartusiak, Piotr Bielak

November 25, 2019



HR EXCELLENCE IN RESEARCH

# Overview

## Concepts

Stateless vs Stateful application

High Availability

Continuous X

Software Updates

## Resource orchestration platforms

Hadoop

YARN

Zookeeper

Swarm

Kubernetes

Openstack

## Useful tools

Helm

Ansible

# Overview

## Concepts

Stateless vs Stateful application

High Availability

Continuous X

Software Updates

## Resource orchestration platforms

Hadoop

YARN

Zookeeper

Swarm

Kubernetes

Openstack

## Useful tools

Helm

Ansible

# Latency times

## Concepts

<https://bit.ly/34jmK47>

### Latency Numbers Every Programmer Should Know

■ 1 ns

■ L1 cache reference: 0.5 ns

■ Branch mispredict: 5 ns

■ L2 cache reference: 7 ns

■ Mutex lock/unlock: 25 ns

■ 100 ns

■ Main memory reference: 100 ns

■ 1 μs

■ Compress 1 KB with Zippy: 3 μs

■ 10 μs

■ Send 1 KB over 1 Gbps network: 10 μs

■ 550 random read (100B/s SSD): 150 μs

■ Read 1 MB sequentially from memory: 250 μs

■ Round trip in same datacenter: 500 μs

■ 1 ms

■ Read 1 MB sequentially from SSD: 1 ms

■ Disk seek: 10 ms

■ Read 1 MB sequentially from disk: 20 ms

■ Packet roundtrip CA to Netherlands: 150 ms

Source: <https://gist.github.com/2841832>

# Latency times

## Concepts

- ▶ L1 cache reference = 0.5 ns
- ▶ Branch mispredict = 5 ns
- ▶ L2 cache reference = 7 ns
- ▶ Mutex lock/unlock = 25 ns
- ▶ Main memory reference = 100 ns
- ▶ Compress 1K bytes with Zippy = 3,000 ns = 3  $\mu$ s
- ▶ Send 2K bytes over 1 Gbps network = 20,000 ns = 20  $\mu$ s
- ▶ SSD random read = 150,000 ns = 150  $\mu$ s
- ▶ Read 1 MB sequentially from memory = 250,000 ns = 250  $\mu$ s
- ▶ Round trip within same datacenter = 500,000 ns = 0.5 ms
- ▶ Read 1 MB sequentially from SSD = 1,000,000 ns = 1 ms
- ▶ Disk seek = 10,000,000 ns = 10 ms
- ▶ Read 1 MB sequentially from disk = 20,000,000 ns = 20 ms
- ▶ Send packet CA  $\rightarrow$  Netherlands  $\rightarrow$  CA = 150,000,000 ns = 150 ms

# Stateless vs Stateful applications

## Concepts

- ▶ How to design apps?
- ▶ How to store user sessions?
- ▶ How to store user data?

# Stateless (1)

## Concepts (Stateless vs Stateful application)

- ▶ Application do not persist any data in their memory
- ▶ Requests are processed one by one
- ▶ Requests are processed independently (without context of previous requests)
- ▶ No session persisted in the app memory

# Stateless (2)

## Concepts (Stateless vs Stateful application)

- ▶ Session can be persisted in persistence solution
  - ▶ Accessible by other replicas of app
  - ▶ For example Redis or MongoDB
  - ▶ Can be scaled and replicated separately
- ▶ There is no need of sticky session usage
- ▶ App can be easily scaled horizontally
- ▶ Development process might be more complex



# Stateful (1)

## Concepts (Stateless vs Stateful application)

- ▶ Application can persist data in memory
- ▶ Requests can be processed in context of previous requests
- ▶ Session is persisted in app memory

# Stateful (2)

## Concepts (Stateless vs Stateful application)

- ▶ We need to use sticky session
- ▶ Scaling and availability is not so trivial
- ▶ Development might be easier

# High Availability

## Concepts (High Availability)

- ▶ characteristic of production level systems
- ▶ should be online (available) as much as possible
- ▶ SLA (Service Level Agreement) - "nines"


Availability %	Downtime per year <sup>[note 1]</sup>	Downtime per month	Downtime per week	Downtime per day
55.55555555% ("nine fives")	162.33 days	13.53 days	74.92 hours	10.67 hours
<b>90% ("one nine")</b>	<b>36.53 days</b>	<b>73.05 hours</b>	<b>16.80 hours</b>	<b>2.40 hours</b>
95% ("one nine five")	18.26 days	36.53 hours	8.40 hours	1.20 hours
97%	10.96 days	21.92 hours	5.04 hours	43.20 minutes
98%	7.31 days	14.61 hours	3.36 hours	28.80 minutes
<b>99% ("two nines")</b>	<b>3.65 days</b>	<b>7.31 hours</b>	<b>1.68 hours</b>	<b>14.40 minutes</b>
99.5% ("two nines five")	1.83 days	3.65 hours	50.40 minutes	7.20 minutes
99.8%	17.53 hours	87.66 minutes	20.16 minutes	2.88 minutes
<b>99.9% ("three nines")</b>	<b>8.77 hours</b>	<b>43.83 minutes</b>	<b>10.08 minutes</b>	<b>1.44 minutes</b>
99.95% ("three nines five")	4.38 hours	21.92 minutes	5.04 minutes	43.20 seconds
<b>99.99% ("four nines")</b>	<b>52.60 minutes</b>	<b>4.38 minutes</b>	<b>1.01 minutes</b>	<b>8.64 seconds</b>
99.995% ("four nines five")	26.30 minutes	2.19 minutes	30.24 seconds	4.32 seconds
<b>99.999% ("five nines")</b>	<b>5.26 minutes</b>	<b>26.30 seconds</b>	<b>6.05 seconds</b>	<b>864.00 milliseconds</b>
<b>99.9999% ("six nines")</b>	<b>31.56 seconds</b>	<b>2.63 seconds</b>	<b>604.80 milliseconds</b>	<b>86.40 milliseconds</b>
<b>99.99999% ("seven nines")</b>	<b>3.16 seconds</b>	<b>262.98 milliseconds</b>	<b>60.48 milliseconds</b>	<b>8.64 milliseconds</b>
<b>99.999999% ("eight nines")</b>	<b>315.58 milliseconds</b>	<b>26.30 milliseconds</b>	<b>6.05 milliseconds</b>	<b>864.00 microseconds</b>
<b>99.9999999% ("nine nines")</b>	<b>31.56 milliseconds</b>	<b>2.63 milliseconds</b>	<b>604.80 microseconds</b>	<b>86.40 microseconds</b>

# HA idea

## Concepts (High Availability)

How to achieve HA?

- ▶ replicate your service
- ▶ route all the traffic through a proxy / load balancer



# HA in real systems

## Concepts (High Availability)



But in reality it's not that easy:

- ▶ implementation of health checks, liveness probes,
- ▶ retry mechanisms,
- ▶ how to automatically scale services?
- ▶ choose scaling criterion, number of fallback replicas,
- ▶ ensure fast application startup,

# Active-active vs active-passive

## Concepts (High Availability)

- ▶ two types of HA: **active-active**, **active-passive**
- ▶ for statefull application:
  - ▶ active-active is hard to maintain → the state of all instances must be synchronized along with processing of requests,
  - ▶ active-passive requires also state synchronization (easier, only one node receives requests),
- ▶ stateless applications - both HA types are easier to apply,



# Continuous X

## Concepts

- ▶ How to cooperate within teams?
- ▶ What we can do to make sure our software works and has high quality
- ▶ How to find out when something was broken?
- ▶ What if two people change dependant functionalities?

# Continuous Integration (1)

## Concepts (Continuous X)

Term comes from Extreme Programming by Kent Beck

- ▶ testing your code is important!,
- ▶ different kind of tests:
  - ▶ unit,
  - ▶ integration,
  - ▶ functional,
  - ▶ regression,
  - ▶ code quality,
  - ▶ security / bugs tests



# Continuous Integration (2)

## Concepts (Continuous X)

- ▶ manual vs automated tests,
- ▶ manual:
  - ▶ not perfect solution, but still used
  - ▶ e.g. done by testers to prepare automated tests
  - ▶ acceptance tests on client side, etc.
- ▶ automated:
  - ▶ utilize testing frameworks (e.g. in Python: unittest, pytest, flake8, in Scala: ScalaTest, Play),
  - ▶ should handle creation of proper testing environments (reproducible builds),
  - ▶ can be ran on developers machine,
  - ▶ after a developer pushes code to remote repository, the CI systems should build the code and run all tests against it (CI pipeline)

# Continuous Deployment (1)

## Concepts (Continuous X)

- ▶ What if continuous integration is not enough for you?
- ▶ What if you want to deploy your software as fast as possible after feature development?
- ▶ How to perform multiple deployments daily? (e.g. Facebook performs thousands of deployments per day)
- ▶ You need to be sure that your production environment is as fresh as it can be?

# Continuous Deployment (2)

## Concepts (Continuous X)

- ▶ Deploy your software as part of your build pipeline (to production or production-like environment)
- ▶ Create push-button deployment possibilities
- ▶ Automate as much as you can
- ▶ CI do not guarantee that your software will be deploy-able, CD does

# Continuous Deployment (3)

## Concepts (Continuous X)

- ▶ Create continuous integration pipeline with all its requirements
- ▶ Extend it to contain executable creation
- ▶ Create automation procedure that allows you to deploy your app/services
- ▶ Utilize appropriate updates procedure (canary etc.)
- ▶ Add deployment as part of build pipeline, or create special build for that

# Automation servers

## Concepts (Continuous X)

### Software development automation servers

- ▶ **Jenkins**
- ▶ **GitLab CI**
- ▶ **TeamCity**
- ▶ **Zuul**
- ▶ **Travis**
- ▶ **CircleCI**

# Software Updates

## Concepts

- ▶ How to perform app updates without downtime?
- ▶ How to perform app updates to part of users
- ▶ How to make sure that app will be working correctly before full roll-out

# Software Update Types

## Concepts (Software Updates)

- ▶ Blue-Green
- ▶ Canary
- ▶ Rolling

# Data schema updates (1)

## Concepts (Software Updates)

- ▶ What about persistence?
- ▶ You need to separate schema from app upgrades
- ▶ On schema changes, you might need to do it in multi-phases



# Data schema updates (2)

## Concepts (Software Updates)

### Field removal:

- ▶ Schema and app requires some field
- ▶ Change schema and app to make field optional
- ▶ Make upgrade
- ▶ Shut down app version that requires that field
- ▶ Change schema and app to make field absent

# Blue-green (1)

## Concepts (Software Updates)

- ▶ Load-balancer of proxy in front of app
- ▶ Two same environments
  - ▶ Green - One environment currently running
  - ▶ Blue - One that is updated
- ▶ Redirect traffic from green to blue after update
- ▶ All ok - blue becomes green, green becomes blue
- ▶ Problem? - redirect traffic back

# Blue-green (2)

## Concepts (Software Updates)

### Update procedure

- ▶ Update blue environment
- ▶ Redirect traffic from green to blue
- ▶ All OK - blue becomes green, green can be again used as blue
- ▶ Problem? - redirect traffic back

# Canary (1)

## Concepts (Software Updates)

- ▶ Roll-out updates only to small part of traffic
- ▶ You can roll-out only to part of your users
- ▶ At Facebook:
  - ▶ Gatekeeper
  - ▶ Roll-out using information about users
  - ▶ By age
  - ▶ By gender
  - ▶ etc.

# Canary (2)

## Concepts (Software Updates)

### Update procedure

- ▶ Create app instance with new features
- ▶ Redirect part of traffic using some predicate to new instance
- ▶ All OK? - continue redirecting the traffic until 100% on new app
- ▶ Problem - roll-back to the old instances

# Rolling (1)

## Concepts (Software Updates)

- ▶  $N$  - number of running instances before update
- ▶ During the update, at most  $N + 1$  running instances
- ▶ At each step of update, replace one old version instance with new one

# Rolling (2)

## Concepts (Software Updates)

### Update procedure

- ▶ Create new version instance
- ▶ All OK - Turn off one old version instance
- ▶ All OK - Repeat until you have only new version instances
- ▶ Problem? - Recreate old instances

# Overview

## Concepts

Stateless vs Stateful application

High Availability

Continuous X

Software Updates

## Resource orchestration platforms

Hadoop

YARN

Zookeeper

Swarm

Kubernetes

Openstack

## Useful tools

Helm

Ansible



# Overview

## Resource orchestration platforms (Hadoop)

- ▶ framework for processing big data
- ▶ not a brand name
- ▶ open source project
- ▶ effective for analytics, sometimes operations

	<b>Operational</b>	<b>Analytical</b>
Latency	1 ms - 100 ms	1 min - 100 min
Concurrency	1000 - 100,000	1 - 10
Access Pattern	Writes and Reads	Reads
Queries	Selective	Unselective
Data Scope	Operational	Retrospective
End User	Customer	Data Scientist
Technology	NoSQL	MapReduce, BSP

# Introduction

## Resource orchestration platforms (Hadoop)

Hadoop's core:

- ▶ HDFS
- ▶ MapReduce
- ▶ YARN
- ▶ Common

# HDFS Overview

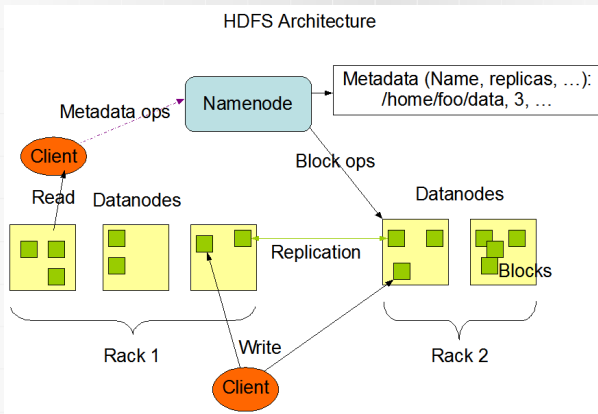
## Resource orchestration platforms (Hadoop)

- ▶ Hadoop Distributed File System
- ▶ heart of Hadoop technology
- ▶ manages how data files are divided and stored across the cluster
- ▶ data is divided into blocks; each server in the cluster contains data from different blocks
- ▶ some built-in redundancy
- ▶ provides file permissions and authentication
- ▶ has a master/slave architecture

# HDFS Architecture

HDFS cluster consists of:

- ▶ single NameNode - master server that manages the file system namespace and regulates access to files
- ▶ a number of DataNodes (usually one per node in the cluster) - manage storage attached to the nodes



# HDFS Operations

## Resource orchestration platforms (Hadoop)

### The NameNode:

- ▶ executes file system namespace
- ▶ opening, closing, and renaming files and directories
- ▶ mapping of blocks to DataNodes

### The DataNodes are responsible for:

- ▶ serving read and write requests from the clients
- ▶ perform block creation, deletion (NameNode instructions)

### Block

- ▶ user data is stored in the files of HDFS; the file in a file system is divided into one or more segments and stored in individual data nodes
- ▶ file segments are called as blocks
- ▶ the minimum amount of data that HDFS can read or write is called a Block
- ▶ default block size is 64MB

# MapReduce

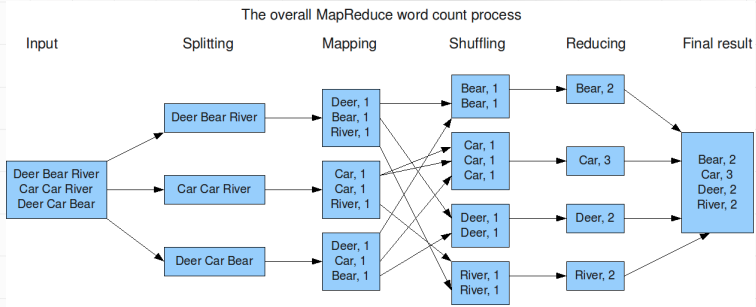
## Resource orchestration platforms (Hadoop)

MapReduce - a method for parallel processing on distributed servers

Stages of MapReduce:

1. map - process the input data from HDFS; input file is passed to the mapper function line by line; creates several small chunks of data
2. shuffle - transfers the map output from Mapper to a Reducer
3. reduce - takes the output of the mapper (intermediate key-value pair); output of the reducer is the final output in HDFS; usually aggregation or summation

# MapReduce example

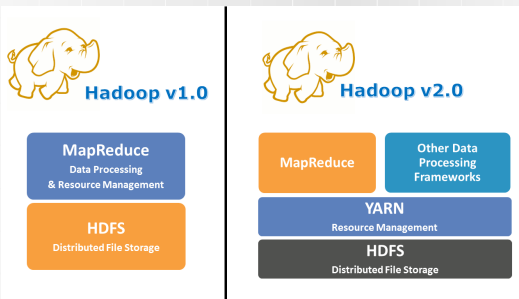


1. the input is first splitting into blocks
2. the splitted input is assined to [key,value] pair in mapper side
3. all the values are ordered in shuffling phase
4. reducer side sorting and grouping is done

# Hadoop YARN

## Resource orchestration platforms (YARN)

- ▶ not to be confused with Javascript's yarn package manager
- ▶ handles processing part in Hadoop
- ▶ YARN = Yet Another Resource Negotiator
- ▶ assigns resources to each application





# YARN components

## Resource orchestration platforms (YARN)

### ► Resource Manager

- communicates with the client;
- tracks resources on the cluster;
- coordinates node manager by assigning job according to requirement;
- allocates memory in form of containers;

### ► Node Manager

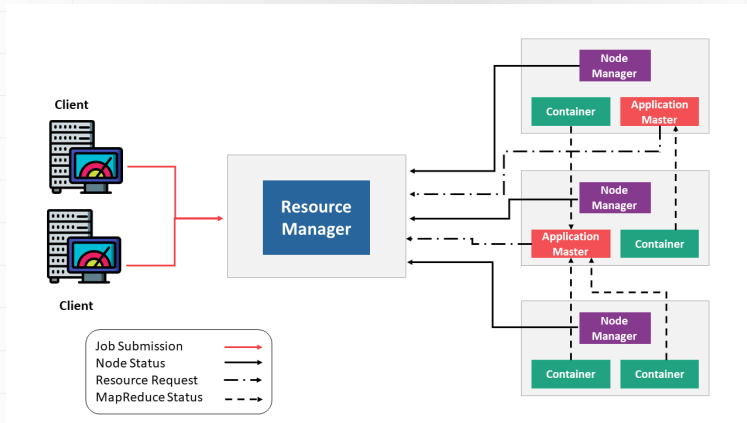
- monitoring of running app;
- offers resources (memory and CPUs) as resource container;
- every datanode → one node manager;
- containers run tasks (including Application Masters);

### ► Application Master

- every job is initialized by Application manager;
- manages the global assignment of compute resources to applications (e.g. memory, CPU, disk, network, etc.)

# YARN diagram

## Resource orchestration platforms (YARN)



# Zookeeper

## Resource orchestration platforms (Zookeeper)


- ▶ not a resource orchestration tool
- ▶ often used along with Hadoop, HBase, Hive, Kafka etc.
- ▶ centralized service for:
  - ▶ maintaining configuration information,
  - ▶ naming,
  - ▶ providing distributed synchronization,
  - ▶ providing group services.



# Docker Swarm

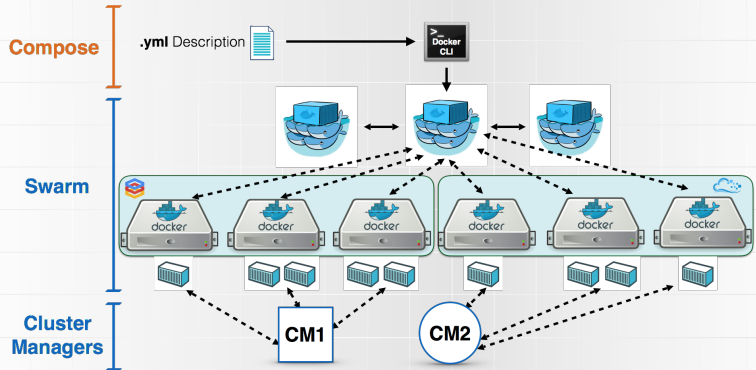
## Resource orchestration platforms (Swarm)

- ▶ container orchestration system created by Docker;
- ▶ allows to distribute Docker containers across multiple nodes;
- ▶ introduces *stacks* and services;
- ▶ resource reservations and limits;
- ▶ container scaling;
- ▶ rolling updates;
- ▶ does not provide advanced networking (overlay);
- ▶ very basic compared to Kubernetes;



# Docker Swarm

## Resource orchestration platforms (Swarm)



# Kubernetes

## Resource orchestration platforms (Kubernetes)

- ▶ Google
- ▶ Based on huge experience on handling distributed environments
- ▶ Open-source
- ▶ huge community

# Kubernetes

## Resource orchestration platforms (Kubernetes)

- ▶ Not so easy to deploy
- ▶ Many dependant services
- ▶ K3S
  - ▶ a simpler and easier in deployment and management
  - ▶ use SQLite, for big traffic might be a bottleneck
- ▶ For local development
  - ▶ Minikube
  - ▶ microk8s

# Architecture

## Resource orchestration platforms (Kubernetes)

Kubernetes has two types of nodes (each running multiple microservices):

- ▶ master (control plane) nodes:

- ▶ API server,
- ▶ Scheduler,
- ▶ Controller Manager,
- ▶ ETCD,

- ▶ worker nodes:

- ▶ Kubelet,
- ▶ Kube-proxy,
- ▶ *Container runtime*\*



# Master nodes components (1)

## Resource orchestration platforms (Kubernetes)

<https://kubernetes.io/docs/concepts/overview/components/>

- ▶ API server (*kube-apiserver*):
  - ▶ exposes the Kubernetes API,
  - ▶ front end for the control plane,
  - ▶ designed to scale horizontally (via load balancers),
- ▶ Scheduler (*kube-scheduler*):
  - ▶ watches newly created pods that have no node assigned,
  - ▶ selects a node for them to run on,
  - ▶ individual and collective resource requirements, hardware/software/policy constraints, affinity and anti-affinity specifications, data locality, inter-workload interference and deadlines,

# Master nodes components (2)

## Resource orchestration platforms (Kubernetes)

<https://kubernetes.io/docs/concepts/overview/components/>

- ▶ CM (*kube-controller-manager*) – runs controllers:
  - ▶ Node Controller – Responsible for noticing and responding when nodes go down.
  - ▶ Replication Controller – Responsible for maintaining the correct number of pods for every replication controller object in the system.
  - ▶ Endpoints Controller – Populates the Endpoints object (that is, joins Services & Pods).
  - ▶ Service Account & Token Controllers – Create default accounts and API access tokens for new namespaces,
- ▶ ETCD:
  - ▶ consistent and highly-available key value store,
  - ▶ used as Kubernetes' backing store for all cluster data,
  - ▶ can be run in standalone or cluster mode,

# Worker nodes components

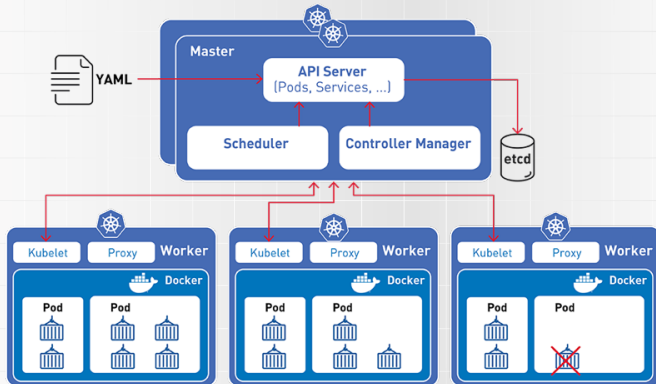
## Resource orchestration platforms (Kubernetes)

<https://kubernetes.io/docs/concepts/overview/components/>

- ▶ kubelet:
  - ▶ daemon running on each worker node,
  - ▶ makes sure that containers are running in a pod,
  - ▶ manages only containers created by Kubernetes,
- ▶ kube-proxy:
  - ▶ network proxy running on each worker node,
  - ▶ maintains network rules on nodes (ingress / egress),
  - ▶ utilizes OS packet filtering layer (if available)
- ▶ Container runtime – supported ones:
  - ▶ Docker, containerd,
  - ▶ cri-o, rktlet,
  - ▶ any implementation of the Kubernetes CRI (Container Runtime Interface).

# Architecture diagram

## Resource orchestration platforms (Kubernetes)



# Kubernetes objects

Resource orchestration platforms (Kubernetes)

kubernetes defines multiple objects that an application can make use of:

- ▶ deployment
- ▶ service
- ▶ daemonset
- ▶ statefulset
- ▶ configmap
- ▶ secret
- ▶ persistentvolume(claim)

# Kubernetes manifests

## Resource orchestration platforms (Kubernetes)

Kubernetes object are created using YAML manifests:

```
1  ---
2  apiVersion: apps/v1
3  kind: Deployment
4  metadata:
5    name: nginx-deployment
6  spec:
7    selector:
8      matchLabels:
9        app: nginx
10   replicas: 1
11   template:
12     metadata:
13       labels:
14         app: nginx
15   spec:
16     volumes:
17     - name: nginx-sample-page
18       configMap:
19         name: nginx-sample-page
20   containers:
21   - name: nginx
22     image: nginx:latest
23     ports:
24     - containerPort: 80
25     volumeMounts:
26     - name: nginx-sample-page
27       mountPath: /usr/share/nginx/html/index.html
28       subPath: index.html
```

# Kubernetes management

## Resource orchestration platforms (Kubernetes)

- ▶ most cluster actions are performed using `kubectl`,
- ▶ it allows to manage the cluster, as well as deploy application manifests,
- ▶ alternatively, for deployment you can use `helm` (we'll discuss it later),

# Openstack

## Resource orchestration platforms (Openstack)


- ▶ most popular open source virtual machine orchestration platform;
- ▶ operates on VMs instead on containers;
- ▶ microservice based architecture;
- ▶ multiple specialized services;





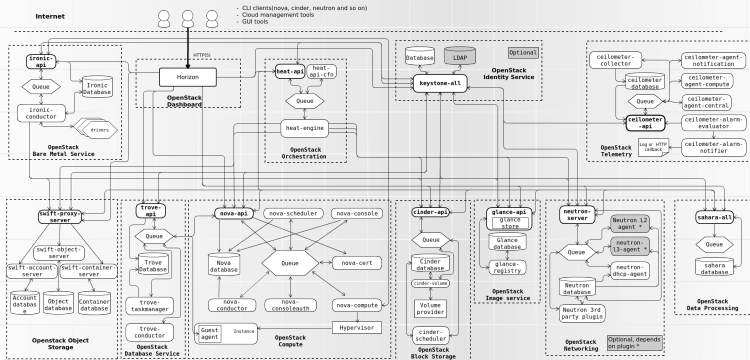
# Architecture (1)

## Resource orchestration platforms (Openstack)



## Architecture (2)

## Resource orchestration platforms (Openstack)



# Services

## Resource orchestration platforms (Openstack)

- ▶ **NOVA** (Compute)
- ▶ **ZUN** (Containers)
- ▶ **QINLING** (Functions)
- ▶ **IRONIC** (Bare Metal Provisioning)
- ▶ **CYBORG** (Hardware accelerators)
- ▶ **SWIFT** (Object store)
- ▶ **CINDER** (Block Storage)
- ▶ **MANILA** (Shared filesystems)
- ▶ **NEUTRON** (Networking)
- ▶ **OCTAVIA** (Load balancer)
- ▶ **DESIGNATE** (DNS)
- ▶ **KEYSTONE** (Identity)
- ▶ **PLACEMENT** (Placement)
- ▶ **GLANCE** (Image)
- ▶ **HEAT** (Orchestration)
- ▶ **MISTRAL** (Workflow)
- ▶ **AODH** (Alarming)
- ▶ **MAGNUM** (Container Orchestration Engine Provisioning)
- ▶ **SAHARA** (Big Data Processing Framework Provisioning)
- ▶ **TROVE** (Database as a Service)
- ▶ **MASAKARI** (Instances High Availability Service)
- ▶ **MURANO** (Application Catalog)
- ▶ **EC2API** (EC2 API proxy)
- ▶ **HORIZON** (Dashboard)
- ▶ and many more...

# Overview

## Concepts

Stateless vs Stateful application

High Availability

Continuous X

Software Updates

## Resource orchestration platforms

Hadoop

YARN

Zookeeper

Swarm

Kubernetes

Openstack

## Useful tools

Helm

Ansible

# Helm

## Useful tools

- ▶ How to manage Kubernetes deployments?
- ▶ There are not variables in Kubernetes manifests
- ▶ What to do if we want to rollback to previous version

# Helm

## Useful tools

- ▶ Package manager for K8S
- ▶ OpenSource, maintained by Google, Microsoft, Bitnami, ...
- ▶ Many ready to use charts in official repository
- ▶ Possibility to use variables during deployment
- ▶ Charts can have dependencies



# Helm

## Useful tools

- ▶ Easily find packages
- ▶ Easily create new packages
- ▶ Can operate on any K8S cluster
- ▶ Query the cluster to see installed packages
- ▶ Update, delete, rollback and check history of installed charts

# Helm

## Useful tools

- ▶ Basically it templates K8S manifests (Go templates)
- ▶ Thanks to that, it is independent from resources that you want to create
- ▶ Create resource YAML, utilize `{{ }}` to inject variables
- ▶ Variables can be provided as a file or during the deployment



# Helm

## Useful tools

### Helm v2

- ▶ requires Tiler
  - ▶ server-side (on cluster) component
  - ▶ manages packages
- ▶ 2-way strategic merge path
- ▶ releases names are global
- ▶ release name is optional, if not provided will be generated

### Helm v3

- ▶ no Tiler
- ▶ 3-way strategic merge path
- ▶ releases names are in namespaces
- ▶ release name is required
- ▶ library charts

# Directory structure

## Useful tools (Helm)

```
1  wordpress/
2    Chart.yaml      # A YAML file containing information
3                    # about the chart
4    LICENSE         # OPTIONAL: A plain text file containing
5                    # the license for the chart
6    README.md       # OPTIONAL: A human-readable README file
7    values.yaml     # The default configuration values
8                    # for this chart
9    values.schema.json # OPTIONAL: A JSON Schema for imposing
10                     # a structure on the values.yaml file
11    charts/         # A directory containing any charts upon
12                     # which this chart depends.
13    crds/           # Custom Resource Definitions
14    templates/      # A directory of templates that,
15                     # when combined with values, will
16                     # generate valid Kubernetes manifest files.
17    templates/NOTES.txt # OPTIONAL: A plain text file containing
18                     # short usage notes
```

# Example template

## Useful tools (Helm)

```
1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: {{ .Release.Name }}-configmap
5  data:
6    myvalue: "Hello World"
```

# Example template

## Useful tools (Helm)

```
1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: {{ .Release.Name }}-configmap
5  data:
6    myvalue: "Hello World"
7    drink: {{ .Values.favoriteDrink }}
```

```
1  helm install --set favoriteDrink=monsterek ./mychart
```

# Helm

## Useful tools

- ▶ How to check if your template will work?
- ▶ dry-run will render template to STD
- ▶ you need to check it manually
- ▶ rendered template can be not accepted by kubernetes

# Introduction

## Useful tools (Ansible)

- ▶ open source automation tool,
- ▶ machine provision,
- ▶ configuration management,
- ▶ application deployment,
- ▶ YAML-based declarative language,
- ▶ agentless (uses SSH / Powershell),



A N S I B L E

# Characteristics

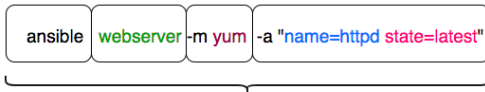
## Useful tools (Ansible)

- ▶ simple & minimalistic (YAML-based language & Jinja templates),
- ▶ consistency (of created environments),
- ▶ security (no dedicated agent, SSH used for connections),
- ▶ reliability (**idempotent** playbooks),

# Command vs playbook

## Useful tools (Ansible)

### AD HOC command



### Ansible Playbook

```
---  
- name: playbook name  
  hosts: webserver  
  tasks:  
    - name: name of the task  
      yum:  
        name: httpd  
        state: latest
```



# Example playbook

## Useful tools (Ansible)

```
1  ---
2  - name: Install nginx
3    hosts: all
4    become: true
5
6    tasks:
7      - name: Add epel-release repo
8        yum:
9          name: epel-release
10         state: present
11
12     - name: Install nginx
13       yum:
14         name: nginx
15         state: present
16
17     - name: Insert Index Page
18       template:
19         src: index.html
20         dest: /usr/share/nginx/html/index.html
21
22     - name: Start NGiNX
23       service:
24         name: nginx
25         state: started
```

# Concepts (1)

## Useful tools (Ansible)

- ▶ playbooks:
  - ▶ define steps to build environments,
  - ▶ can be divided into multiple files (readability, reusability),
  - ▶ roles, vars, group\_vars etc.
- ▶ modules:
  - ▶ define actual actions executed by Ansible,
  - ▶ examples given on previous slide,
  - ▶ standalone,
  - ▶ can be written in most scripting languages (Python, Bash, Perl, Ruby etc.),
  - ▶ should follow **idempotent** rule,

# Concepts (2)

## Useful tools (Ansible)

- ▶ inventory file:
  - ▶ description of nodes that can be accessed by Ansible,
  - ▶ INI or YAML format,
  - ▶ IP addresses or hostnames,
  - ▶ when necessary SSH keys and users can be provided,
  - ▶ nodes can be assigned to groups,

# Inventory file example

## Useful tools (Ansible)

```
1  # Consolidation of all groups
2  [hosts:children]
3  web-servers
4  offsite
5  onsite
6  backup-servers
7
8  [web-servers]
9  server1 ansible_host=192.168.0.1 ansible_port=1600
10 server2 ansible_host=192.168.0.2 ansible_port=1800
11
12 [offsite]
13 server3 ansible_host=10.160.40.1 ansible_port=22 ansible_user=root
14 192.168.6.1
15
16 # You can make groups of groups
17 [offsite:children]
18 backup-servers
19
20 [onsite]
21 server5 ansible_host=10.150.70.1 ansible_ssh_pass=password
22
23 [backup-servers]
24 foo.example.com
```

# Large Scale Data Processing

## Lecture 3 – Data processing stack

dr inż. Tomasz Kajdanowicz, Roman Bartusiak, Piotr Bielak

November 25, 2019